

Fast Bilateral Filtering

Sylvain Paris, Adobe
sparis@adobe.com

- Manipulating the texture in a photo is a central operation in many tasks.
- Let's look at a few examples...

Photographic Style Transfer

[Bae 06]



input

Photographic Style Transfer

[Bae 06]



output

Tone Mapping

[Durand 02]



HDR input

Tone Mapping

[Durand 02]



Cartoon Rendition

[Winnemöller 06]



Cartoon Rendition


[Winnemöller 06]



Naïve Approach: Gaussian Blur



input



BLUR



*smoothed
(structure, large scale)*



HALOS




*residual
(texture, small scale)*

Gaussian Convolution

Impact of Blur and Halos

- If the decomposition introduces blur and halos, the final result is corrupted.


Sample manipulation:
increasing texture
(residual $\times 3$)



Bilateral Filter: no Blur, no Halos



input



smoothed
(structure, large scale)



residual
(texture, small scale)

edge-preserving: Bilateral Filter

input





increasing texture
with Gaussian convolution

H A L O S



increasing texture
with bilateral filter
N O H A L O S

Traditional Denoising versus Computational Photography

Edge-preserving filtering introduced for denoising.

- Denoising: decompose into signal + noise
 - Throw away noise
 - Small kernels
- Computational photography: decompose into base + detail
 - Detail is valuable
 - Large kernels
 - ↳ Bilateral filter [Aurich 95, Smith 97, Tomasi 98]

Objective of bilateral filtering

Smooth texture


Preserve edges

Illustration a 1D Image

- 1D image = line of pixels



- Better visualized as a plot




Definition

Gaussian blur

$$I_{\mathbf{p}}^b = \sum_{\mathbf{q} \in \mathcal{S}} \underbrace{G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|)}_{\text{space}} I_{\mathbf{q}}$$

- only spatial distance, intensity ignored




Bilateral filter

[Aurich 95, Smith 97, Tomasi 98]

$$I_{\mathbf{p}}^{\text{bf}} = \underbrace{\frac{1}{W_{\mathbf{p}}^{\text{bf}}}}_{\text{normalization}} \sum_{\mathbf{q} \in \mathcal{S}} \underbrace{G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|)}_{\text{space}} \underbrace{G_{\sigma_r}(|I_{\mathbf{p}} - I_{\mathbf{q}}|)}_{\text{range}} I_{\mathbf{q}}$$


normalization

- spatial and range distances
- weights sum to 1



Example on a Real Image

- Kernels can have complex, spatially varying shapes.



Bilateral Filter is Expensive

- Brute-force computation is slow (several minutes)
 - Two nested for loops:
for each pixel, look at all pixels
 - Non-linear, depends on image content
⇒ no FFT, no pre-computation...
- Fast approximations exist [Durand 02, Weiss 06]
 - Significant **loss of accuracy**
 - **No formal understanding** of accuracy versus speed

Today


- We will reformulate the bilateral filter
 - Link with **linear filtering**
 - **Fast** and **accurate** algorithm

Questions ?


Outline

- Reformulation of the BF
- Fast algorithm to compute the BF
- Practical implementation
- Application and extension
 - Photographic style transfer
 - Bilateral grid


Bilateral Filter on 1D Signal



BF



Our Strategy




Reformulate the bilateral filter

- More complex space:
 - Homogeneous intensity
 - Higher-dimensional space
- Simpler expression: mainly a convolution
 - ↳ Leads to a fast algorithm

Link with Linear Filtering

1. Handling the Division



sum of
weights

$$I_{\mathbf{p}}^{\text{bf}} = \frac{1}{W_{\mathbf{p}}^{\text{bf}}} \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(|I_{\mathbf{p}} - I_{\mathbf{q}}|) I_{\mathbf{q}}$$

Handling the division with a **projective space**.

Formalization: Handling the Division

$$I_{\mathbf{p}}^{\text{bf}} = \frac{1}{W_{\mathbf{p}}^{\text{bf}}} \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_{\text{s}}}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_{\text{r}}}(|I_{\mathbf{p}} - I_{\mathbf{q}}|) I_{\mathbf{q}}$$
$$W_{\mathbf{p}}^{\text{bf}} = \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_{\text{s}}}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_{\text{r}}}(|I_{\mathbf{p}} - I_{\mathbf{q}}|)$$

- Normalizing factor as homogeneous coordinate
 - Multiply both sides by $W_{\mathbf{p}}^{\text{bf}}$

$$\begin{pmatrix} W_{\mathbf{p}}^{\text{bf}} I_{\mathbf{p}}^{\text{bf}} \\ W_{\mathbf{p}}^{\text{bf}} \end{pmatrix} = \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_{\text{s}}}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_{\text{r}}}(|I_{\mathbf{p}} - I_{\mathbf{q}}|) \begin{pmatrix} I_{\mathbf{q}} \\ 1 \end{pmatrix}$$

Formalization: Handling the Division

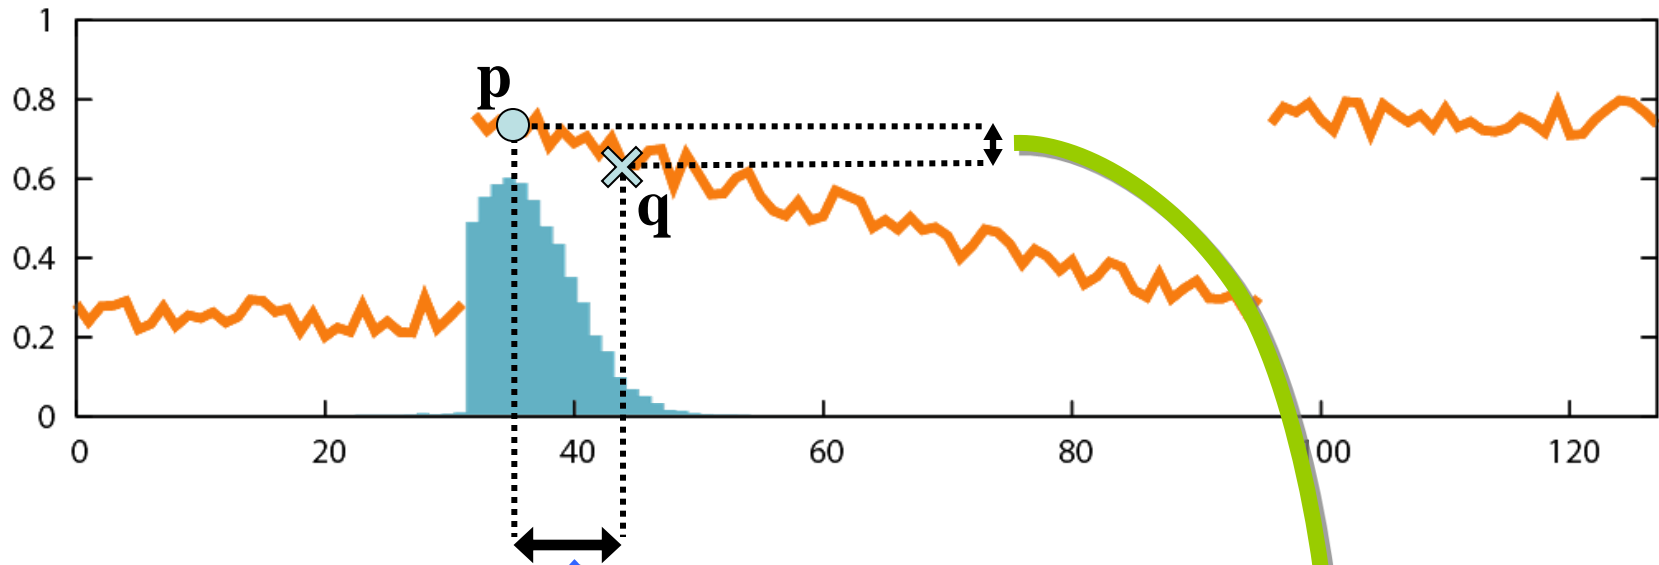
$$\begin{pmatrix} W_{\mathbf{p}}^{\text{bf}} & I_{\mathbf{p}}^{\text{bf}} \\ & W_{\mathbf{p}}^{\text{bf}} \end{pmatrix} = \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(|I_{\mathbf{p}} - I_{\mathbf{q}}|) \begin{pmatrix} W_{\mathbf{q}} & I_{\mathbf{q}} \\ & W_{\mathbf{q}} \end{pmatrix} \text{ with } W_{\mathbf{q}}=1$$

- Similar to homogeneous coordinates in projective space
- Division delayed until the end

Questions ?

Link with Linear Filtering

2. Introducing a Convolution




$$\begin{pmatrix} W_{\mathbf{p}}^{\text{bf}} & I_{\mathbf{p}}^{\text{bf}} \\ W_{\mathbf{p}}^{\text{bf}} & \end{pmatrix} = \sum_{\mathbf{q} \in \mathcal{S}} \underbrace{G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|)}_{\text{space}} \underbrace{G_{\sigma_r}(|I_{\mathbf{p}} - I_{\mathbf{q}}|)}_{\text{range}} \begin{pmatrix} W_{\mathbf{q}} & I_{\mathbf{q}} \\ W_{\mathbf{q}} & \end{pmatrix}$$

2D Gaussian

Link with Linear Filtering

2. Introducing a Convolution




$$\begin{pmatrix} W_{\mathbf{p}}^{\text{bf}} & I_{\mathbf{p}}^{\text{bf}} \\ W_{\mathbf{p}}^{\text{bf}} \end{pmatrix} = \sum_{\mathbf{q} \in \mathcal{S}} \boxed{G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(|I_{\mathbf{p}} - I_{\mathbf{q}}|)} \begin{pmatrix} W_{\mathbf{q}} & I_{\mathbf{q}} \\ W_{\mathbf{q}} \end{pmatrix}$$

space x range

Corresponds to a 3D Gaussian on a 2D image.
Result appeared previously in [Barash 02].

Link with Linear Filtering

2. Introducing a Convolution




sum all values

$$\begin{pmatrix} W_{\mathbf{p}}^{\text{bf}} & I_{\mathbf{p}}^{\text{bf}} \\ W_{\mathbf{p}}^{\text{bf}} & \end{pmatrix} = \boxed{\sum_{(\mathbf{q}, \zeta) \in \mathcal{S} \times \mathcal{R}} \sum} \quad \text{space-range Gaussian} \quad \begin{pmatrix} W_{\mathbf{q}} & I_{\mathbf{q}} \\ W_{\mathbf{q}} & \end{pmatrix}$$

sum all values multiplied by kernel \Rightarrow convolution

Link with Linear Filtering

2. Introducing a Convolution




result of the convolution

$$\begin{pmatrix} W_{\mathbf{p}}^{\text{bf}} & I_{\mathbf{p}}^{\text{bf}} \\ W_{\mathbf{p}}^{\text{bf}} \end{pmatrix} = \sum_{(\mathbf{q}, \zeta) \in \mathcal{S} \times \mathcal{R}} \sum \text{space-range Gaussian} \begin{pmatrix} W_{\mathbf{q}} & I_{\mathbf{q}} \\ W_{\mathbf{q}} \end{pmatrix}$$


Link with Linear Filtering

2. Introducing a Convolution




result of the convolution


$$\begin{pmatrix} W_{\mathbf{p}}^{\text{bf}} & I_{\mathbf{p}}^{\text{bf}} \\ W_{\mathbf{p}}^{\text{bf}} \end{pmatrix} = \sum_{(\mathbf{q}, \zeta) \in \mathcal{S} \times \mathcal{R}} \text{space-range Gaussian} \begin{pmatrix} W_{\mathbf{q}} & I_{\mathbf{q}} \\ W_{\mathbf{q}} \end{pmatrix}$$




higher dimensional functions




Gaussian convolution



division



slicing



Reformulation: Summary

linear: $(w^{\text{bf}} i^{\text{bf}}, w^{\text{bf}}) = g_{\sigma_s, \sigma_r} \otimes (wi, w)$

nonlinear: $I_{\mathbf{p}}^{\text{bf}} = \frac{w^{\text{bf}}(\mathbf{p}, I_{\mathbf{p}}) i^{\text{bf}}(\mathbf{p}, I_{\mathbf{p}})}{w^{\text{bf}}(\mathbf{p}, I_{\mathbf{p}})}$

1. Convolution in higher dimension

- expensive but well understood (linear, FFT, etc)

2. Division and slicing


- nonlinear but simple and pixel-wise

Exact reformulation

Questions ?

Outline

- Reformulation of the BF
- Fast algorithm to compute the BF
- Practical implementation
- Application and extension
 - Photographic style transfer
 - Bilateral grid



higher dimensional functions

Recap:

- simple operations
- complex space



w_i


w

Gaussian convolution

division

slicing





higher dimensional functions

Strategy:
downsampled
convolution

w_i

w



DOWNSAMPLE

Gaussian convolution

UPSAMPLE


division

slicing



Sampling Theorem


- Sampling a signal at a least twice its smallest wavelength is enough.



Not enough

Sampling Theorem


- Sampling a signal at a least twice its smallest wavelength is enough.



Not enough


Sampling Theorem

- Sampling a signal at a least twice its smallest wavelength is enough.



Enough

Signal processing analysis




higher dimensional functions




Low-pass filter


Gaussian convolution




division



slicing



“Safe”
downsampling



higher dimensional functions




DOWNSAMPLE

Gaussian convolution




UPSAMPLE

division



slicing




Fast Convolution by Downsampling

- Downsampling cuts frequencies above Nyquist limit (half the sampling rate)
 - Less data to process
 - But introduces error
- Evaluation of the approximation
- Efficient implementation


Accuracy versus Running Time

≈1 second instead of **several minutes**

- Finer sampling increases accuracy.
- More precise than previous work.




Digital
photograph
1200 × 1600



Brute-force bilateral
filter takes over 10
minutes.

Visual Results



1200 × 1600

- Comparison with previous work [Durand 02]
 - running time = 1s for both techniques

input



exact BF




our result





prev. work



difference
with exact
computation
(intensities in [0:1])





More on Accuracy and Running Times



Kernel Size

- **Larger kernels are faster** because we downsample more.
↳ Useful for photography.



Digital
photograph
 1200×1600

Brute-force bilateral
filter takes over 10
minutes.

Questions ?

Outline

- Reformulation of the BF
- Fast algorithm to compute the BF
- Practical implementation
- Application and extension
 - Photographic style transfer
 - Bilateral grid

Efficient Implementation

- Never build the full resolution 3D space
 - Bin pixels on the fly
 - Interpolate on the fly
- Separable Gaussian kernel
- 5-tap approximation

Sampling Rate

- 1 sample every sigma
 - Kernel parameters are all equal to 1 sample
 - 5-tap approximation is sufficient

$$1 - 4 - 6 - 4 - 1$$

FAST BILATERAL FILTER

input: image I
Gaussian parameters σ_s and σ_r
sampling rates s_s and s_r

output: filtered image I^b

1. Initialize all w_{\downarrow} i_{\downarrow} and w_{\downarrow} values to 0. **Initialize the 3D grid to 0.**

FAST BILATERAL FILTER

input: image I
Gaussian parameters σ_s and σ_r
sampling rates s_s and s_r

output: filtered image I^b

1. Initialize all w_{\downarrow} , i_{\downarrow} and w_{\uparrow} values to 0. **Initialize the 3D grid to 0.**
2. Compute the minimum intensity value:

Useful later to save space. $I_{\min} \leftarrow \min_{(X,Y) \in \mathcal{S}} I(X,Y)$

3. For each pixel $(X, Y) \in \mathcal{S}$ with an intensity $I(X, Y) \in \mathcal{R}$ **Look at each pixel.**

(a) Compute the homogeneous vector (wi, w) :

Create the data. $(wi, w) \leftarrow (I(X, Y), 1)$

3. For each pixel $(X, Y) \in \mathcal{S}$ with an intensity $I(X, Y) \in \mathcal{R}$ **Look at each pixel.**

(a) Compute the homogeneous vector (wi, w) :

Create the data. $(wi, w) \leftarrow (I(X, Y), 1)$

(b) Compute the downsampled coordinates (with $[\cdot]$ the rounding operator)

Compute the grid location. $(x, y, \zeta) \leftarrow \left(\left[\frac{X}{s_s} \right], \left[\frac{Y}{s_s} \right], \left[\frac{I(X, Y) - I_{\min}}{s_r} \right] \right)$

3. For each pixel $(X, Y) \in \mathcal{S}$ with an intensity $I(X, Y) \in \mathcal{R}$ **Look at each pixel.**

(a) Compute the homogeneous vector (wi, w) :

Create the data. $(wi, w) \leftarrow (I(X, Y), 1)$

(b) Compute the downsampled coordinates (with $[\cdot]$ the rounding operator)

Compute the grid location. $(x, y, \zeta) \leftarrow \left(\left[\frac{X}{s_s} \right], \left[\frac{Y}{s_s} \right], \left[\frac{I(X, Y) - I_{\min}}{s_r} \right] \right)$

(c) Update the downsampled $\mathcal{S} \times \mathcal{R}$ space

Update the grid.
$$\begin{pmatrix} w_{\downarrow} i_{\downarrow}(x, y, \zeta) \\ w_{\downarrow}(x, y, \zeta) \end{pmatrix} \leftarrow \begin{pmatrix} w_{\downarrow} i_{\downarrow}(x, y, \zeta) \\ w_{\downarrow}(x, y, \zeta) \end{pmatrix} + \begin{pmatrix} wi \\ w \end{pmatrix}$$

4. Convolve $(w_{\downarrow} i_{\downarrow}, w_{\downarrow})$ with a 3D Gaussian g whose parameters are σ_s/s_s and σ_r/s_r
- $$(w_{\downarrow}^b i_{\downarrow}^b, w_{\downarrow}^b) \leftarrow (w_{\downarrow} i_{\downarrow}, w_{\downarrow}) \otimes g$$

**Convolve with 1 – 4 – 6 – 4 – 1 along each axis.
3 for loops needed, one for each axis.**

**In 3D, 15 samples (3 times 5) considered
instead of 125 (5^3) for a full convolution.
Same result!**

5. For each pixel $(X, Y) \in \mathcal{S}$ with an intensity $I(X, Y) \in \mathcal{R}$ **Look at each pixel.**

(a) Tri-linearly interpolate the functions $w_{\downarrow}^b i_{\downarrow}^b$ and w_{\downarrow}^b to obtain $W^b I^b$ and W^b :

$$\begin{aligned} W^b I^b(X, Y) &\leftarrow \text{interpolate} \left(w_{\downarrow}^b i_{\downarrow}^b, \frac{X}{s_s}, \frac{Y}{s_s}, \frac{I(X, Y)}{s_r} \right) \\ W^b(X, Y) &\leftarrow \text{interpolate} \left(w_{\downarrow}^b, \frac{X}{s_s}, \frac{Y}{s_s}, \frac{I(X, Y)}{s_r} \right) \end{aligned}$$

5. For each pixel $(X, Y) \in \mathcal{S}$ with an intensity $I(X, Y) \in \mathcal{R}$ **Look at each pixel.**

(a) Tri-linearly interpolate the functions $w_{\downarrow}^b i_{\downarrow}^b$ and w_{\downarrow}^b to obtain $W^b I^b$ and W^b :

$$W^b I^b(X, Y) \leftarrow \text{interpolate} \left(w_{\downarrow}^b i_{\downarrow}^b, \frac{X}{s_s}, \frac{Y}{s_s}, \frac{I(X, Y)}{s_r} \right)$$

$$W^b(X, Y) \leftarrow \text{interpolate} \left(w_{\downarrow}^b, \frac{X}{s_s}, \frac{Y}{s_s}, \frac{I(X, Y)}{s_r} \right)$$

(b) Normalize the result

$$I^b(X, Y) \leftarrow \frac{W^b I^b(X, Y)}{W^b(X, Y)}$$

Comments

- Every sample is processed even if empty
 - The grid is coarse and fairly dense in 3D.
 - e.g. parameters (16,0.1): 256 pixels for 10 bins
convolution spans 5 bins
at least 50% occupancy
 - Simple data structure, simple sweep \Rightarrow fast
 - More sophisticated approaches needed for higher dimensional cases
[Adams et al. 09,10] [Gastal and Oliveira 11,12]

Complexity

- There is no nested loops over the whole set of samples
 - At most: “for each sample, for 5 samples”

$$\mathcal{O} \left(|\mathcal{S}| + \frac{|\mathcal{S}|}{s_s^2} \frac{|\mathcal{R}|}{s_r} \right)$$

Creation + slicing
“for each pixel”



Convolution
“for each grid sample”



References

- Short version:** A Fast Approximation of the Bilateral Filter using a Signal Processing Approach. *Sylvain Paris and Frédo Durand*. ECCV 2006
- Long version:** A Fast Approximation of the Bilateral Filter using a Signal Processing Approach. *Sylvain Paris and Frédo Durand*. International Journal of Computer Vision, 2009
- Different presentation, more applications:** Real-time Edge-Aware Image Processing with the Bilateral Grid. *Jiawen Chen, Sylvain Paris, and Frédo Durand*. SIGGRAPH 2007
- Survey:** Bilateral Filtering: Theory and Applications. *Sylvain Paris, Pierre Kornprobst, Jack Tumblin, and Frédo Durand*. Foundations and Trends in Computer Graphics and Vision, 2009
- Data, code, etc:** <http://people.csail.mit.edu/sparis/bf/>